

Machine Learning — Repeat Exam

1	2	3	4	5	6	7	8	9	10	11	Σ
4	5	6	4	6	5	4	7	4	2	7	54

Do not write anything above this line

Name:

Student ID:

Signature:

- Only write on the sheets given to you by supervisors. If you need more paper, ask the supervisors.
- Pages 16-18 can be used as scratch paper.
- All sheets (including scratch paper) have to be returned at the end.
- **Do not unstaple the sheets!**
- Wherever answer boxes are provided, please write your answers in them.
- Please write your student ID (*Matrikelnummer*) on every sheet you hand in.
- **Only use a black or a blue pen (no pencils, red or green pens!).**
- You are allowed to use your A4 sheet of handwritten notes (two sides). **No other materials (e.g. books, cell phones, calculators) are allowed!**
- Exam duration - 120 minutes.
- This exam consists of 18 pages, 11 problems. You can earn 54 points.

Probability distributions

For your reference, we provide the following probability distribution.

- Univariate normal distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

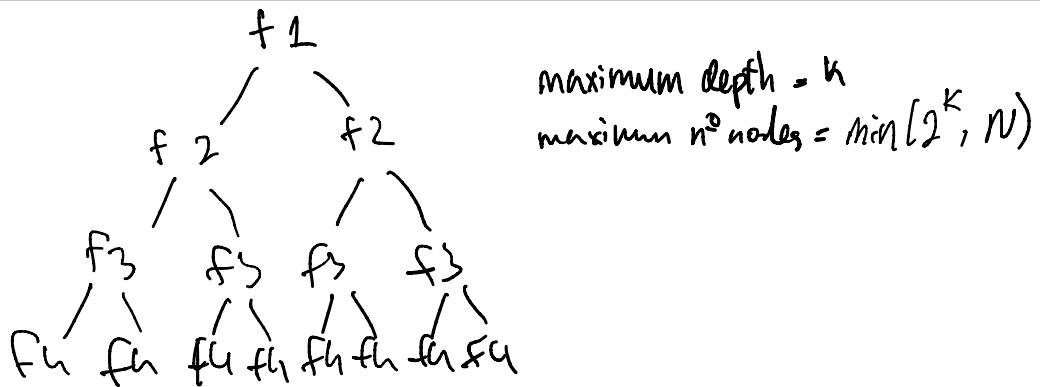
- Bernoulli distribution

$$\text{Bern}(x|\theta) = \theta^x(1-\theta)^{1-x}$$

Decision Trees

Problem 1 [(2+2)=4 points] Assume you want to build a decision tree. Your data set consists of N samples, each with k features ($k \leq N$).

- a) If the features are binary, what is the maximum possible number of leaf nodes and the maximum depth of your decision tree?



- b) If the features are continuous, what is the maximum possible number of leaf nodes and the maximum depth of your decision tree?

Max. n^o leaf nodes : N
Max. depth : N

Regression

Problem 2 [(1+4)=5 points] We want to perform regression on a dataset consisting of N samples $x_i \in \mathbb{R}^D$ with corresponding targets $y_i \in \mathbb{R}$ (represented compactly as $\mathbf{X} \in \mathbb{R}^{N \times D}$ and $\mathbf{y} \in \mathbb{R}^N$).

Assume that we have fitted an L_2 -regularized linear regression model and obtained the optimal weight vector $\mathbf{w}^* \in \mathbb{R}^D$ as

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

Note that there is no bias term.

Now, assume that we obtained a new data matrix \mathbf{X}_{new} by scaling all samples by the same positive factor $a \in (0, \infty)$. That is, $\mathbf{X}_{new} = a\mathbf{X}$ (and respectively $\mathbf{x}_i^{new} = ax_i$).

- a) Find the weight vector \mathbf{w}_{new} that will produce the same predictions on \mathbf{X}_{new} as \mathbf{w}^* produces on \mathbf{X} .

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{w}^{*T} \mathbf{x}_i = \mathbf{W}_{new}^T \mathbf{x}_i^{new} \Rightarrow \mathbf{w}^{*T} \mathbf{x}_i = \mathbf{W}_{new}^T \alpha \mathbf{x}_i ; \quad \mathbf{W}_{new} = \frac{\mathbf{W}^*}{\alpha} \Rightarrow \mathbf{W}_{new} = \frac{\mathbf{W}^*}{\alpha}$$

- b) Find the regularization factor $\lambda_{new} \in \mathbb{R}$, such that the solution \mathbf{w}_{new}^* of the new L_2 -regularized linear regression problem

$$\mathbf{w}_{new}^* = \arg \min_{\mathbf{w}} \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \mathbf{x}_i^{new} - y_i)^2 + \frac{\lambda_{new}}{2} \mathbf{w}^T \mathbf{w}$$

will produce the same predictions on \mathbf{X}_{new} as \mathbf{w}^* produces on \mathbf{X} .

Provide a mathematical justification for your answer.

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

$$\begin{aligned} \mathbf{w}_{new}^* &= (\mathbf{X}_{new}^T \mathbf{X}_{new} + \lambda_{new} \mathbf{I})^{-1} \mathbf{X}_{new}^T \mathbf{y} \\ &= \alpha (\alpha^2 \mathbf{X}^T \mathbf{X} + \lambda_{new} \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$

by setting $\lambda_{new} = \alpha^2 \lambda$, we get

$$= \alpha (\alpha^2 \mathbf{X}^T \mathbf{X} + \alpha^2 \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \frac{1}{\alpha} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \frac{1}{\alpha} \mathbf{w}^*$$

which according to our answer in a, will produce the same predictions

on \mathbf{Y}_{new} as \mathbf{w}^* does on \mathbf{X} , as desired

Classification

Problem 3 [(1+2+3)=6 points] We would like to perform binary classification on multivariate binary data. That is, the data points $\mathbf{x}_i \in \{0, 1\}^D$ are binary vectors of length D , and each sample belongs to one of two classes $y_i \in \{1, 2\}$.

Consider the following generative classification model. We place a categorical prior on y

$$p(y = 1) = \pi_1 \quad p(y = 2) = \pi_2.$$

The class-conditional distributions are products of independent Bernoulli distributions

$$p(\mathbf{x} | y = 1, \boldsymbol{\alpha}) = \prod_{j=1}^D \text{Bern}(x_j | \alpha_j),$$

$$p(\mathbf{x} | y = 2, \boldsymbol{\beta}) = \prod_{j=1}^D \text{Bern}(x_j | \beta_j),$$

where $\boldsymbol{\alpha} \in [0, 1]^D$ and $\boldsymbol{\beta} \in [0, 1]^D$ are the respective parameter vectors for both classes. That is, each component x_j is distributed as $x_j \sim \text{Bern}(\alpha_j)$ if $y = 1$ or $x_j \sim \text{Bern}(\beta_j)$ if $y = 2$.

- a) Write down the expression for the posterior distribution $p(y | \mathbf{x}, \boldsymbol{\alpha}, \boldsymbol{\beta}, \boldsymbol{\pi})$.

$$\begin{aligned} p(y | \mathbf{x}, \boldsymbol{\alpha}, \boldsymbol{\beta}, \boldsymbol{\pi}) &= \frac{p(\mathbf{x} | y, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y | \boldsymbol{\pi})}{p(\mathbf{x} | \boldsymbol{\alpha}, \boldsymbol{\beta}, \boldsymbol{\pi})} \\ &= \frac{p(\mathbf{x} | y=1, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=1 | \boldsymbol{\pi})}{p(\mathbf{x} | y=1, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=1 | \boldsymbol{\pi}) + p(\mathbf{x} | y=2, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=2 | \boldsymbol{\pi})} \end{aligned}$$

- b) Assume that $D = 3$, $\boldsymbol{\alpha} = [1/3, 1/3, 3/4]$, $\boldsymbol{\beta} = [2/3, 1/2, 1/2]$, $\pi_1 = 1/3$ and $\pi_2 = 2/3$.

Write down a data point $\mathbf{x}_1 \in \{0, 1\}^3$ that will be classified as class 1 by our model. Additionally, compute the posterior probability $p(y = 1 | \mathbf{x}_1, \boldsymbol{\alpha}, \boldsymbol{\beta}, \boldsymbol{\pi})$.

$\mathbf{x}_1 = (0, 0, 1)$. The probability that \mathbf{x}_1 belongs to class 1 is

$$\begin{aligned} p(y=1 | \mathbf{x}_1, \boldsymbol{\alpha}, \boldsymbol{\beta}, \boldsymbol{\pi}) &= \frac{p(\mathbf{x}_1 | y=1, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=1 | \boldsymbol{\pi})}{p(\mathbf{x}_1 | y=1, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=1 | \boldsymbol{\pi}) + p(\mathbf{x}_1 | y=2, \boldsymbol{\alpha}, \boldsymbol{\beta}) \cdot p(y=2 | \boldsymbol{\pi})} \\ &\leq \frac{(1-\alpha_1) \cdot (1-\alpha_2) \cdot \alpha_3 \cdot \pi_1}{(1-\alpha_1) \cdot (1-\alpha_2) \cdot \alpha_3 \cdot \pi_1 + (1-\beta_1) \cdot (1-\beta_2) \cdot \beta_3 \cdot \pi_2} \\ &= \frac{2}{3} \end{aligned}$$

- c) Consider the case when $D = 2$, $\pi_1 = \pi_2 = 1/2$, and $\boldsymbol{\alpha} \in [0, 1]^2$ and $\boldsymbol{\beta} \in [0, 1]^2$ are known and fixed. Show that the resulting classification rule can be represented as a linear function of \mathbf{x} . That is, find $\mathbf{w} \in \mathbb{R}^2$ and $b \in \mathbb{R}$, such that

$$\{\mathbf{x} \in \{0, 1\}^2 : \mathbf{w}^T \mathbf{x} + b > 0\} = \{\mathbf{x} \in \{0, 1\}^2 : p(y = 1 | \mathbf{x}) > p(y = 2 | \mathbf{x})\}$$

$$\begin{aligned}
 & p(y=1|x) > p(y=2|x) \\
 \Leftrightarrow & p(x|y=1)p(y=1) > p(x|y=2)p(y=2) \\
 \Leftrightarrow & p(x|y=1) > p(x|y=2) \\
 \Leftrightarrow & \log p(x|y=1) > \log p(x|y=2) \\
 \Leftrightarrow & \log p(x_1|y=1) + \log p(x_2|y=1) > \log p(x_1|y=2) + \log p(x_2|y=2) \\
 \Leftrightarrow & x_1 \log \alpha_1 + (1-x_1) \log(1-\alpha_1) + x_2 \log \beta_2 + (1-x_2) \log(1-\beta_2) \\
 & - x_2 \log \beta_1 - (1-x_1) \log(1-\beta_1) - x_1 \log \alpha_2 - (1-x_2) \log(1-\alpha_2) > 0 \\
 \Leftrightarrow & (\log \alpha_1 - \log(1-\alpha_1) - \log \beta_1 + \log(1-\beta_1))x_1 \\
 & + (\log \beta_2 - \log(1-\beta_2) - \log \alpha_2 + \log(1-\alpha_2))x_2 \quad \} \in \mathbb{R}^2 \\
 & + (\log(1-\alpha_1) + \log(1-\alpha_2) - \log(1-\beta_1) - \log(1-\beta_2)) > 0 \quad b \in \mathbb{R}
 \end{aligned}$$

Kernels

Problem 4 [(4)=4 points] Prove or disprove whether the following operations on sets $A, B \subseteq \mathcal{X}$, where \mathcal{X} is a finite set, define a valid kernel.

- a) $k(A, B) = |A \times B|$, where $A \times B = \{(a, b) : a \in A, b \in B\}$ denotes the cartesian product and $|S|$ denotes the cardinality of set S , i.e. the number of elements in S .

Yes, using cardinality as a feature map $|A \times B| = |A| \cdot |B|$ obviously defines a kernel

b) $k(A, B) = |A \cap B|$

Since the set X is finite we can denote its members as x_i , with $i \in \{1, 2, \dots, |X|\}$ and define the membership vector a_i with

$$a_i = \begin{cases} 1 & \text{if } x_i \in A \\ 0 & \text{otherwise} \end{cases}$$

Using this as a feature map we can write $|A \cap B| = a^T b$

c) $k(A, B) = |A \cup B|$

Looking at the determinant of a Gram matrix defined by two objects $A, B \subseteq X$ we have

$$\det \begin{pmatrix} |A| & |A \cup B| \\ |A \cup B| & |B| \end{pmatrix} = |A||B| - |A \cup B|^2$$

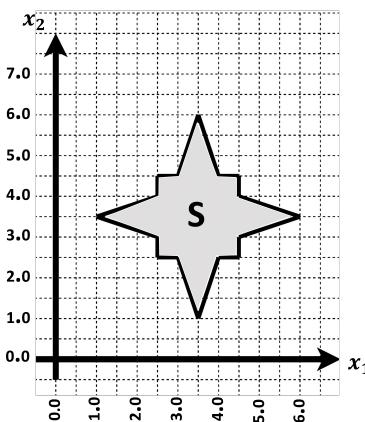
which is negative if $A \neq B$. Hence the Gram matrix is not positive semidefinite and according to Mercer's theorem $|A \cup B|$ is not a kernel

Optimization

Problem 5 [(1+3+2)=6 points] Let f be the following convex function on \mathbb{R}^2 :

$$f(x_1, x_2) = e^{x_1+x_2} - 5 \cdot \log(x_2)$$

a) Consider the following shaded region S in \mathbb{R}^2 . Is this region convex? Why?



No

Because we can choose two points in S where the line connecting the points does not completely reside in S

- b) Find the maximizer (x_1^*, x_2^*) of f over the shaded region S . For your computations, you can pick values from the following table. Justify your answer.

$e^{4.5} = 90.017$	$e^{5.0} = 148.41$	$e^{5.5} = 244.69$	$e^{6.5} = 665.14$
$e^{7.0} = 1096.63$	$e^{7.5} = 1808.04$	$e^{8.0} = 2980.95$	$e^{8.5} = 4914.76$
$e^{9.0} = 8103.08$	$e^{9.5} = 13359.726$	$e^{10.0} = 22026.46$	$e^{10.5} = 36315.50$
$\log(1.0) = 0$	$\log(2.5) = 0.9162$	$\log(3.0) = 1.0986$	$\log(3.5) = 1.2527$
$\log(4.0) = 1.3862$	$\log(4.5) = 1.5040$	$\log(5.0) = 1.6094$	$\log(6.0) = 1.7917$

We need to compute the value of f only on the four corners of the convex hull. Afterwards, we pick the corner at which the value of f is larger.

$$f(1, 3.5) = 83.7533 \quad f(3.5, 1.0) = 90.01$$

$$f(6, 3.5) = 12353.46 \quad f(3.5, 6.0) = 13750.76$$

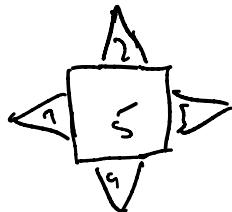
$$\max \rightarrow (6, 3.5), 12353.46$$

.

- c) Assume that we are given an algorithm $\text{ConvOpt}(f, \mathcal{X})$ that takes as input a convex function f and any convex region \mathcal{X} , and returns the minimum of f over \mathcal{X} .

Using the ConvOpt algorithm, how would you find the global minimum of f over the shaded region S ?

Partition the shaded regions to the following 5 convex regions



SVM

Problem 6 [(5)=5 points] Given the data points

$$\mathbf{x}_1 = (1, 1, 0, 1)^T \quad \mathbf{x}_2 = (1, 1, 1, 0)^T \quad \mathbf{x}_3 = (0, 1, 1, 1)^T \quad \mathbf{x}_4 = (0, 0, 1, 1)^T$$

Prove or disprove whether the following combinations of labels \mathbf{y} and dual variables $\boldsymbol{\alpha}$ are the optimal solutions of a soft-margin SVM with $C = 1$.

a) $\mathbf{y} = (-1, -1, 1, 1)^T, \boldsymbol{\alpha} = (0.6, 0.6, 1, 0)^T$

b) $\mathbf{y} = (-1, -1, 1, 1)^T, \boldsymbol{\alpha} = (\frac{2}{3}, \frac{2}{3}, \frac{4}{3}, 0)^T$

c) $\mathbf{y} = (-1, 1, -1, 1)^T, \boldsymbol{\alpha} = (1, 1, 1, 1)^T$

a) No, $\sum_i d_i y_i \neq 0$

b) No, $d_3 = 4/3 > C = 1$

c) This solution is feasible, so we need to check optimality. The dual problem of the soft margin SVM is

$$g(\boldsymbol{\alpha}) = \sum_i d_i - \frac{1}{2} \sum_i \sum_j y_i y_j d_i d_j \mathbf{x}_i^T \mathbf{x}_j$$

To check whether this is an optimal solution we first look at the gradient,

$$(\nabla_{\boldsymbol{\alpha}} g(\boldsymbol{\alpha}))_i = 1 - \sum_j y_j y_j d_j \mathbf{x}_i^T \mathbf{x}_j$$

Deep Learning

Problem 7 [(2+2)=4 points] You are trying to solve a regression task and you want to choose between two approaches:

1. A simple linear regression model.
2. A feed forward neural network $f_{\mathbf{W}}(\mathbf{x})$ with L hidden layers, where each hidden layer $l \in \{1, \dots, L\}$ has a weight matrix $\mathbf{W}_l \in \mathbb{R}^{D \times D}$ and a ReLU activation function. The output layer has a weight matrix $\mathbf{W}_{L+1} \in \mathbb{R}^{D \times 1}$ and no activation function.

In both models, there are no bias terms.

Your dataset \mathcal{D} contains data points with nonnegative features \mathbf{x}_i and the target y_i is continuous:

$$\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^N, \quad \mathbf{x}_i \in \mathbb{R}_{\geq 0}^D, \quad y_i \in \mathbb{R}$$

Let $\mathbf{w}_{LS}^* \in \mathbb{R}^D$ be the optimal weights for the linear regression model corresponding to a global minimum of the following least squares optimization problem:

$$\mathbf{w}_{LS}^* = \arg \min_{\mathbf{w} \in \mathbb{R}^D} \mathcal{L}_{LS}(\mathbf{w}) = \arg \min_{\mathbf{w} \in \mathbb{R}^D} \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \mathbf{x}_i - y_i)^2$$

Let $\mathbf{W}_{NN}^* = \{\mathbf{W}_1^*, \dots, \mathbf{W}_{L+1}^*\}$ be the optimal weights for the neural network corresponding to a global minimum of the following optimization problem:

$$\mathbf{W}_{NN}^* = \arg \min_{\mathbf{W}} \mathcal{L}_{NN}(\mathbf{W}) = \arg \min_{\mathbf{W}} \frac{1}{2} \sum_{i=1}^N (f_{\mathbf{W}}(\mathbf{x}_i) - y_i)^2$$

- a) Assume that the optimal \mathbf{W}_{NN}^* you obtain are non-negative.

What will be the relation ($<, \leq, =, \geq, >$) between the neural network loss $\mathcal{L}_{NN}(\mathbf{W}_{NN}^*)$ and the linear regression loss $\mathcal{L}_{LS}(\mathbf{w}_{LS}^*)$? Provide a mathematical argument to justify your answer.

Any non-negative \mathbf{x} and any non-negative \mathbf{W} it holds $\text{ReLU}(\mathbf{x} \cdot \mathbf{W}) = \mathbf{x} \cdot \mathbf{W}$

Therefore, since our data points have non-negative features \mathbf{x}_i and the optimal weights \mathbf{W}_{NN}^* are non-negative, every ReLU layer is equivalent to a linear layer when plugging in the optimal weights.

$$\begin{aligned} f_{\mathbf{W}_{NN}^*}(\mathbf{x}_i) &= \text{ReLU}(\text{ReLU}(\text{ReLU}(\mathbf{x}_i^T \mathbf{W}_1^*) \mathbf{W}_2^*) \dots \mathbf{W}_L^*) \mathbf{W}_{L+1}^* \\ &= \mathbf{x}_i^T \mathbf{W}_1^* \mathbf{W}_2^* \dots \mathbf{W}_{L+1}^* \\ &= \mathbf{x}_i^T \mathbf{W}_{NN}^* \end{aligned}$$

$$\mathcal{L}_{NN}(\mathbf{W}_{NN}^*) = \mathcal{L}_{LS}(\mathbf{w}_{LS}^*) \quad \mathbf{w}_{NN}^* = \mathbf{w}_{LS}^*$$

- b) In contrast to (a), now assume that the optimal weights w_{LS}^* you obtain are non-negative. What will be the relation ($<$, \leq , $=$, \geq , $>$) between the linear regression loss $\mathcal{L}_{LS}(w_{LS}^*)$ and the neural network loss $\mathcal{L}_{NN}(W_{NN}^*)$? Provide a mathematical argument to justify your answer.

As stated in (a) linear regression is a special case of the above n.n., for any weights w_{LS} you can find weights W_{NN} that produce the same output. That is, everything that can be learned with a linear regression can be learned equally well with a neural network.

However, the reverse direction doesn't hold, since in principle neural networks can learn more complicated functions compared to linear regression.

$\mathcal{L}_{NN}(W_{NN}) \leq \mathcal{L}_{LS}(w_{LS}^*)$ since the nn can potentially find a better fit for the data (taking advantage of non-linearity)

Dimensionality Reduction

Problem 8 [(3+2+2)=7 points] You are given $N = 4$ data points: $\{\mathbf{x}_i\}_{i=1}^4, \mathbf{x}_i \in \mathbb{R}^3$, represented with the matrix $\mathbf{X} \in \mathbb{R}^{4 \times 3}$.

$$\mathbf{X} = \begin{bmatrix} 4 & 3 & 2 \\ 2 & 1 & -2 \\ 4 & -1 & 2 \\ -2 & 1 & 2 \end{bmatrix}$$

Hint: In this task the results of all (final and intermediate) computations happen to be integers.

- a) Perform principal component analysis (PCA) of the data \mathbf{X} , i.e. find the principal components and their associated variances in the transformed coordinate system. Show your work.

Mean = $\begin{bmatrix} 2 & 1 & 1 \end{bmatrix}$
 1) Enter the data $\rightarrow \begin{bmatrix} 2 & 2 & 1 \\ 0 & 0 & -3 \\ -2 & -2 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 2 & 0 & -4 \\ 2 & 0 & 2 & 0 \\ 1 & -3 & 1 & 1 \end{bmatrix} \begin{bmatrix} 2 & 2 & 1 \\ 0 & 0 & -3 \\ -2 & -2 & 1 \end{bmatrix}$

Covariance matrix $\rightarrow \sum \mathbf{x}_c = \frac{1}{N} \mathbf{x}_c^\top \mathbf{x}_c = \begin{bmatrix} 6 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 3 \end{bmatrix}$

$\Lambda = \sum \mathbf{x}_c$ and $T = I_3$ $\sum \mathbf{x}_c = T \cdot \Lambda \cdot T^\top$

- b) Project the data to two dimensions, i.e. write down the transformed data matrix $\mathbf{Y} \in \mathbb{R}^{4 \times 2}$

using the top-2 principal components you computed in (a). What fraction of variance of \mathbf{X} is preserved by \mathbf{Y} ?

$$\mathbf{T}_{\text{trunc}} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{Y} = \mathbf{X}^T \mathbf{T}_{\text{trunc}} = \begin{bmatrix} 2 & 1 \\ 0 & -3 \\ 2 & 1 \\ -4 & 1 \end{bmatrix}$$

We preserve $\frac{6+3}{6+2+3} = \frac{9}{11}$ of the variance

- c) Let $\mathbf{x}_5 \in \mathbb{R}^3$ be a new data point. Specify the vector \mathbf{x}_5 such that performing PCA on the data including the new data point $\{\mathbf{x}_i\}_{i=1}^5$ leads to exactly the same principal components as in (a).

Let $\mathbf{x}_5 = \bar{\mathbf{x}}$ (the new data point equals the mean before including \mathbf{x}_5 to the dataset. Therefore, the new mean including \mathbf{x}_5 is equal to the old mean

$$\mathbf{X}_c = \mathbf{X} - \bar{\mathbf{x}} = \begin{bmatrix} 2 & 2 & 1 \\ 0 & 0 & -3 \\ 2 & -2 & 1 \\ -4 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

which leads to the same $\sum \mathbf{x}_c$ as in (a) up to a difference in the multiplicative constant, in (a) we had $\frac{1}{4} \mathbf{x}_c^T \mathbf{x}_c$ and here we have $\frac{1}{5} \mathbf{x}_c^T \mathbf{x}_c$, while this difference leads to different eigenvalues, the eigenvectors and thus the principal component stay the same

Clustering

Problem 9 [(4)=4 points] Let μ_1, \dots, μ_K be the centroids computed by the K -means algorithm. Prove that the set \mathcal{X}_j of all points in \mathbb{R}^D assigned during inference to the cluster j is a convex set.

$$\mathcal{X}_j := \{x \in \mathbb{R}^D : x \text{ would be assigned to centroid } \mu_j \text{ by } K\text{-means}\}$$

Hint: start by thinking about the case with $K = 2$.

--

Problem 10 [(2)=2 points] Given three 1-dimensional Gaussian distributions $\mathcal{N}(\mu_i, \sigma_i^2)$ with parameters

$$\begin{array}{lll} \mu_1 = 1, & \mu_2 = -1, & \mu_3 = 0, \\ \sigma_1 = 1, & \sigma_2 = 0.5, & \sigma_3 = 2.5 \end{array}$$

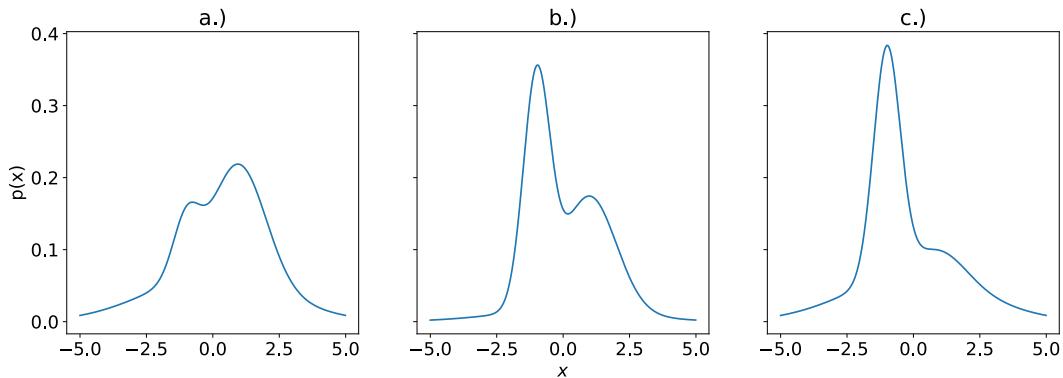
and three different vectors of mixing coefficients π defining categorical cluster priors.

Match the value of π in each row of the following table with one of the probability density functions

$$p(x) = \sum_{i=1}^3 \pi_i \mathcal{N}(x | \mu_i, \Sigma_i)$$

of the resulting GMM showed below. Fill in the last column of the table, no argumentation required.

	π_1	π_2	π_3	PDF (a, b or c)
case 1	0.111...	0.444...	0.444...	c
case 2	0.444...	0.111...	0.444...	a
case 3	0.444...	0.444...	0.111...	b



Variational Inference

Problem 11 [(3+1+1+2)=7 points] Consider the following latent variable probabilistic model

$$\begin{aligned} p(z) &= \mathcal{N}(z | 0, 1) \\ p(x | z) &= \mathcal{N}(x | z, 1) \end{aligned}$$

We want to approximate the posterior distribution $p(z | x)$ using the following variational family

$$\mathcal{Q} = \{\mathcal{N}(z | \mu, 1) \text{ for } \mu \in \mathbb{R}\}$$

that includes all normal distributions with unit variance.

Questions (a), (b), (c) and (d) are all concerning this setup.

Hint: Variance of $p(z | x)$ is equal to 0.5.

- a) Write down the closed-form expression for ELBO $\mathcal{L}(q)$ and simplify it. You can ignore all the terms constant in μ .

- b) Find the optimal variational distribution $q^* \in \mathcal{Q}$ that maximizes the ELBO

$$q^* = \arg \max_{q \in \mathcal{Q}} \mathcal{L}(q)$$

i.e. find the mean μ^* of the optimal variational distribution q^* .

c) Assume that the optimal q^* (i.e., the optimal μ^*) from question (b) is given. Which of the following statements is true?

- (1) $\text{KL}(q(z | \mu^*) \| p(z | x)) < 0$
- (2) $\text{KL}(q(z | \mu^*) \| p(z | x)) = 0$
- (3) $\text{KL}(q(z | \mu^*) \| p(z | x)) > 0$

Justify your answer.

d) For each of the conditions (1), (2), (3) from question (c) above, provide a parametric variational family \mathcal{Q}_i , such that the optimal q_i^* from each family would fulfill the respective condition, or explain why it's impossible.

That is, provide \mathcal{Q}_1 , such that for $q_1^* = \arg \max_{q \in \mathcal{Q}_1} \mathcal{L}(q)$ we have $\text{KL}(q_1^*(z) \| p(z | x)) < 0$, for $q_2^* = \arg \max_{q \in \mathcal{Q}_2} \mathcal{L}(q)$ we have $\text{KL}(q_2^*(z) \| p(z | x)) = 0$, and for $q_3^* = \arg \max_{q \in \mathcal{Q}_3} \mathcal{L}(q)$ we have $\text{KL}(q_3^*(z) \| p(z | x)) > 0$.





